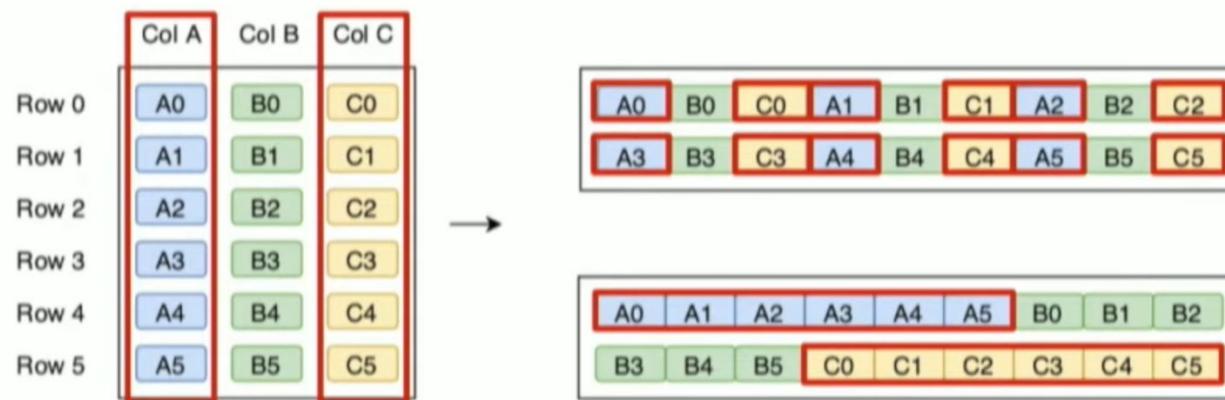


INTRO TO Apache Iceberg

Recap Parquet

Row-wise vs Columnar



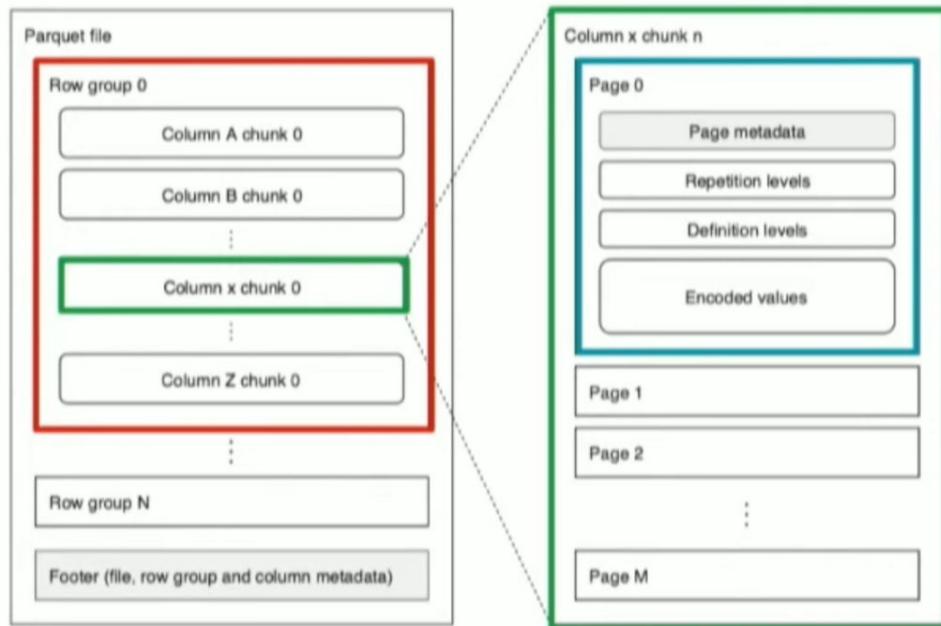
	Col A	Col B	Col C
Row 0	A0	B0	C0
Row 1	A1	B1	C1
Row 2	A2	B2	C2
Row 3	A3	B3	C3
Row 4	A4	B4	C4
Row 5	A5	B5	C5

→

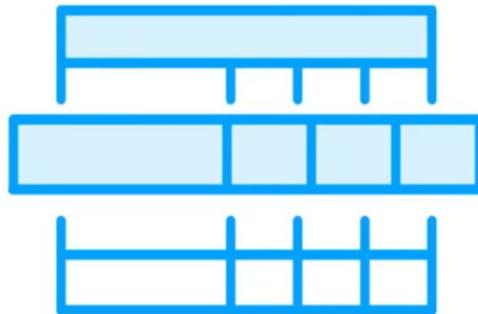
A0	A1	A2	B0	B1	B2	C0	C1	C2
A3	A4	A5	B3	B4	B5	C3	C4	C5

Parquet: data organization

- Data organization
 - Row-groups (*default 128MB*)
 - Column chunks
 - Pages (*default 1MB*)
 - Metadata
 - Min
 - Max
 - Count
 - Rep/def levels
 - Encoded values



How Parquet “Knows” Which Row Group to Scan?



- **Parquet contains metadata**
 - Data about data
- **Min and max values**
- **Footer**
 - Format version
 - Schema information
 - Column metadata

Parquet Storage

	Column 1	Column 2	Column 3	Column 4	Column 5
	Product	Customer	Country	Date	Sales Amount
Row group 1	Ball	John Doe	USA	2023-01-01	100
	T-Shirt	John Doe	USA	2023-01-02	200
Row group 2	Socks	Maria Adams	UK	2023-01-01	300
	Socks	Antonio Grant	USA	2023-01-03	100
Row group 3	T-Shirt	Maria Adams	UK	2023-01-02	500
	Socks	John Doe	USA	2023-01-05	200

Projection and Predicate(s)

Projection = SELECT

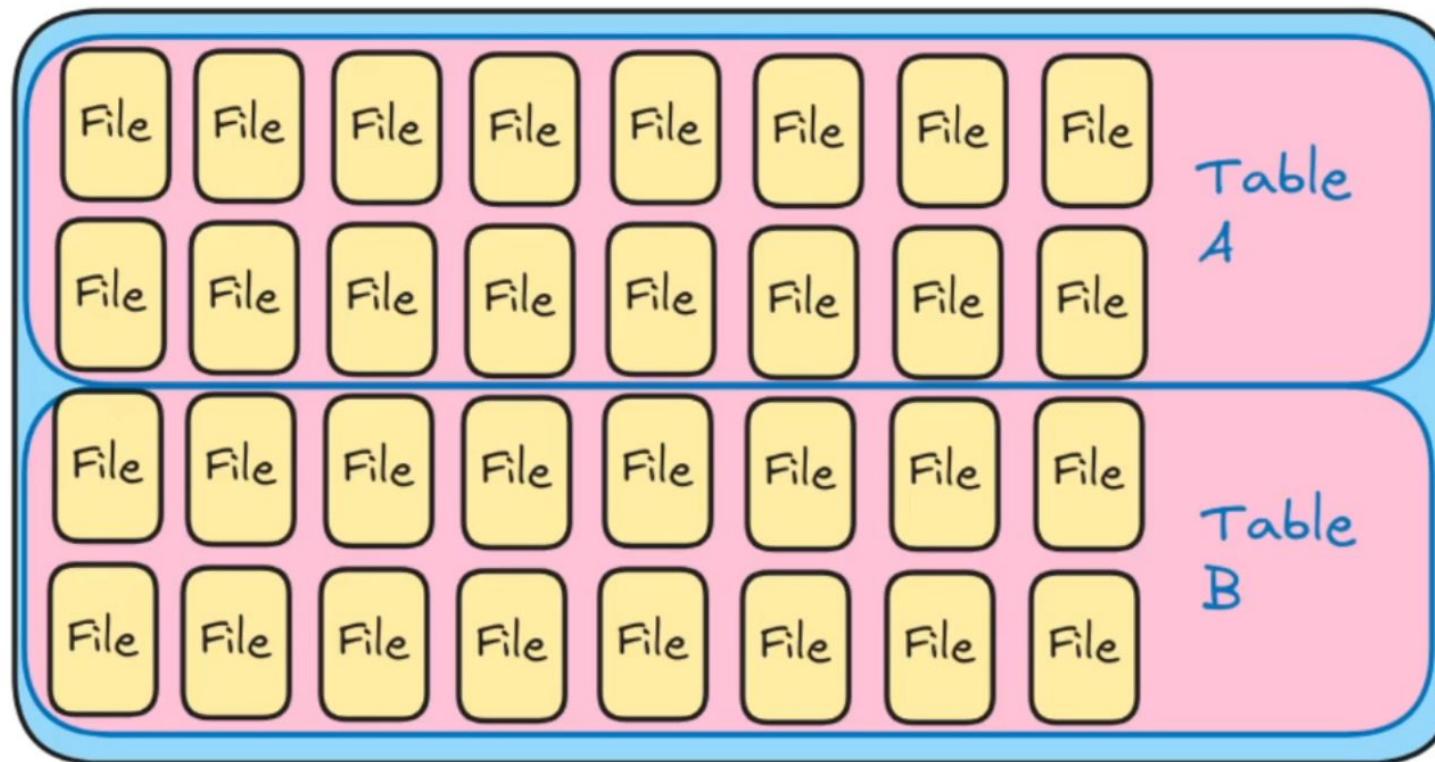
Predicate(s) = WHERE

	Column 1	Column 2	Column 3	Column 4	Column 5
	Product	Customer	Country	Date	Sales
Row group 1	Ball	John Doe	USA	2023-01-01	100
	T-Shirt	John Doe	USA	2023-01-02	200
Row group 2	Socks	The engine will skip scanning these records!			
	Socks	UK	USA	2023-01-02	200
Row group 3	T-Shirt	Maria Adams	UK	2023-01-02	500
	Socks	John Doe	USA	2023-01-05	200



What is a Table Format?

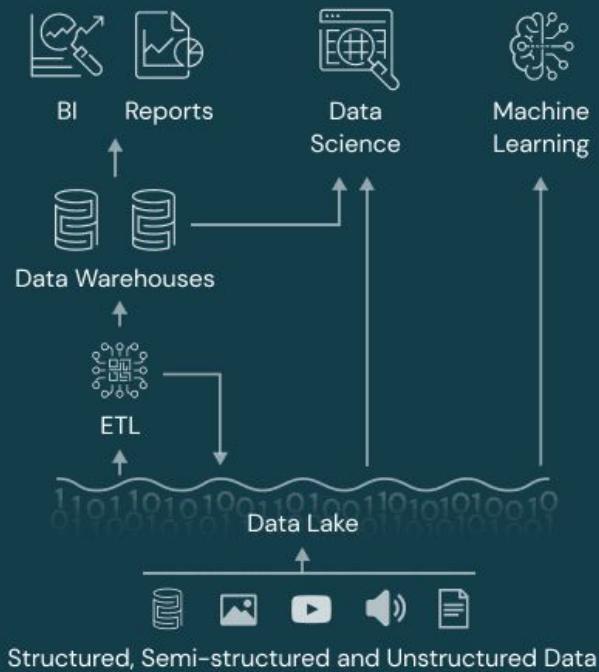
An abstraction to allow multiple files on the lake to be seen as a single table.



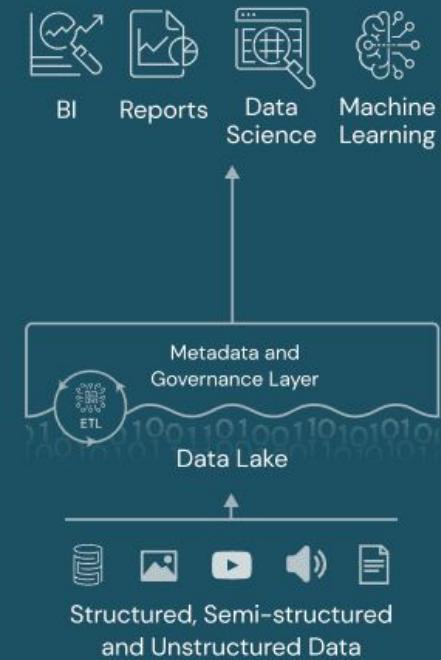
Data Warehouse



Data Lake



Data Lakehouse





Analytics

Storage
Engine



Compute
engine



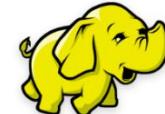
Catalog



Table formats

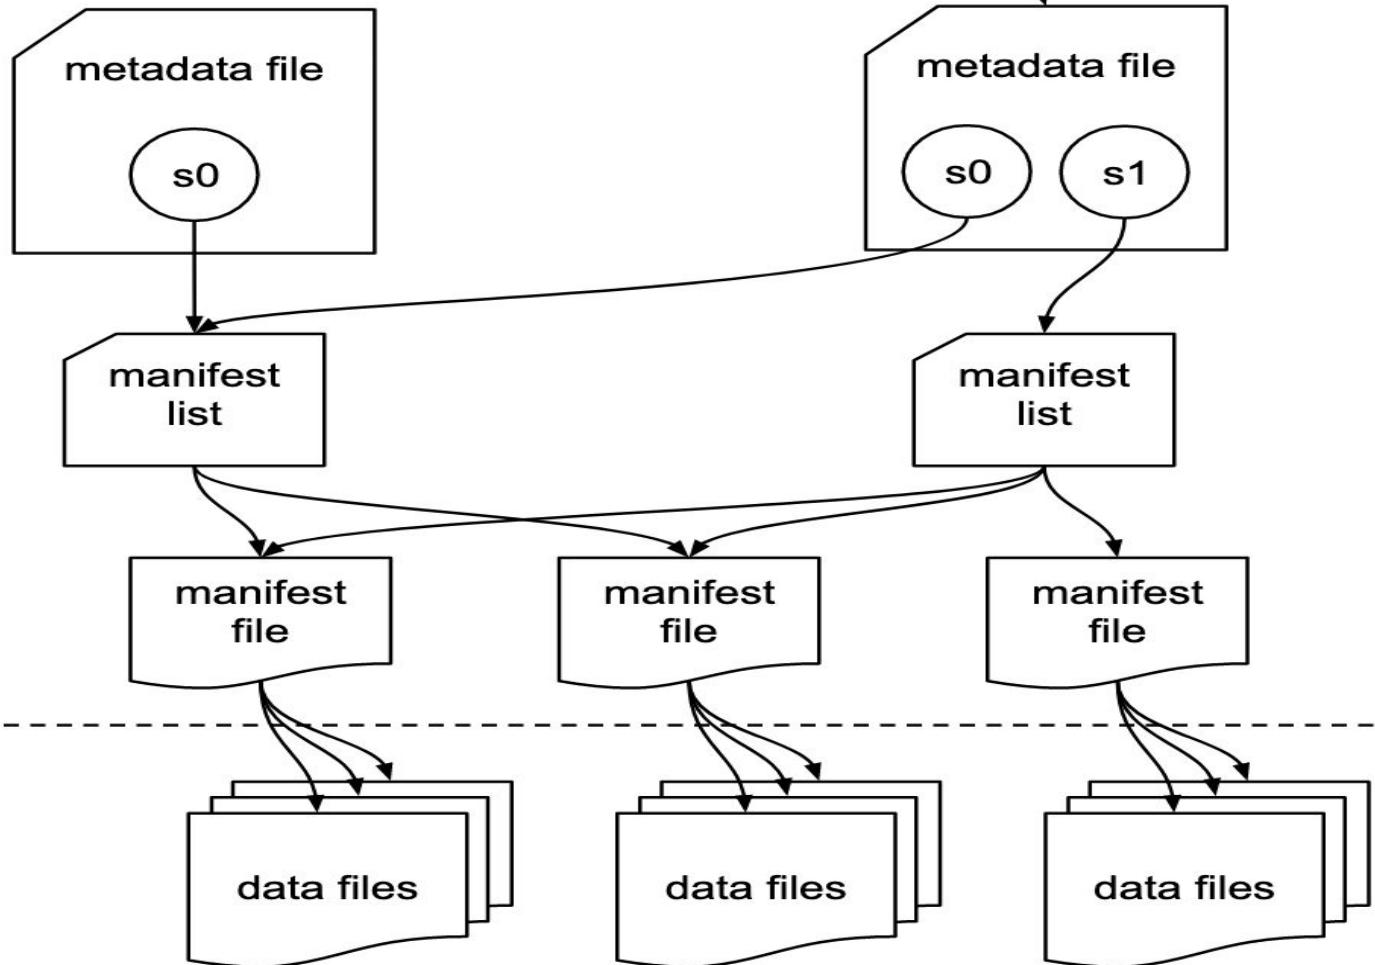


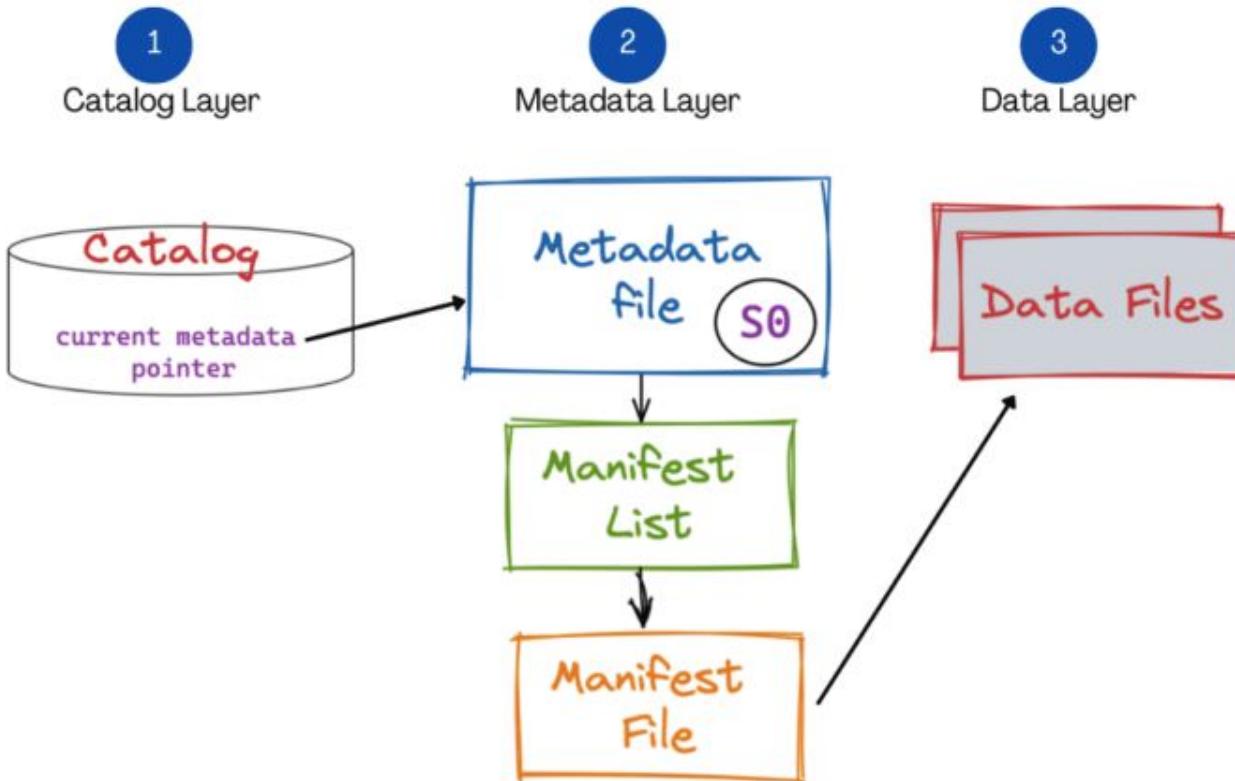
File formats



Lake Storage

Metadata layer





Iceberg Table Metadata File (metadata.json)

```
{  
  "format-version": 2,  
  "table-uuid": "b3d6d8b8-62f1-4fa6-9a8d-08a9f7d4aa31",  
  "location": "s3://my-bucket/warehouse/db/sales",  
  "current-snapshot-id": 3055729220031123456,  
  "snapshots": [  
    {  
      "snapshot-id": 3055729220031123456,  
      "timestamp-ms": 1736940000123,  
      "operation": "append",  
      "manifest-list": ".../metadata/snap-3055729220031123456-1.avro"  
    }  
,  
  "schemas": [  
    { "schema-id": 0, "fields": [ ... ] }  
  ],  
  "partition-specs": [  
    { "spec-id": 0, "fields": [ { "transform": "day(order_ts)" } ] }  
  ]  
}
```

Iceberg Manifest List (Avro)

Manifest List (decoded view) – one record per manifest file

```
[  
  {  
    "manifest_path": ".../metadata/0d3c2b2f-...-m0.avro",  
    "content": "DATA",  
    "partition_spec_id": 0,  
    "sequence_number": 17,  
    "added_snapshot_id": 3055729220031123456,  
    "added_files_count": 2,  
    "existing_files_count": 0,  
    "deleted_files_count": 0,  
    "partitions": [  
      { "lower_bound": "2026-01-15", "upper_bound": "2026-01-16" }  
    ]  
  }  
]
```

Iceberg Manifest File (Avro)

Manifest File (decoded view) – entries point to data/delete files

```
[  
  {  
    "status": "ADDED",  
    "snapshot_id": 3055729220031123456,  
    "sequence_number": 17,  
    "data_file": {  
      "content": "DATA",  
      "file_path": ".../data/order_day=2026-01-15/00000-0-....parquet",  
      "file_format": "PARQUET",  
      "partition": { "order_day": "2026-01-15" },  
      "record_count": 60000,  
      "file_size_in_bytes": 134217728,  
      "lower_bounds": { "order_id": "100000" },  
      "upper_bounds": { "order_id": "159999" }  
    }  
  },  
  {  
    "status": "ADDED",  
    "data_file": {  
      "file_path": ".../data/order_day=2026-01-16/00001-0-....parquet",  
      "record_count": 60000  
    }  
  }  
]
```

```
CREATE TABLE aircraft (
    tail_number varchar(15),
    description varchar(150),
    class varchar(50),
    year integer
)
WITH
    (type = 'iceberg');
```



metadata layer



data layer

[Exit Full Screen](#)

Objects (1) [Info](#)[Copy S3 URI](#)[Copy URL](#)[Download](#)[Open](#)[Delete](#)[Actions ▾](#)[Create folder](#)[Upload](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

 [Find objects by prefix](#) [**1**](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class	
<input type="checkbox"/>	metadata/	Folder	-	-	-	



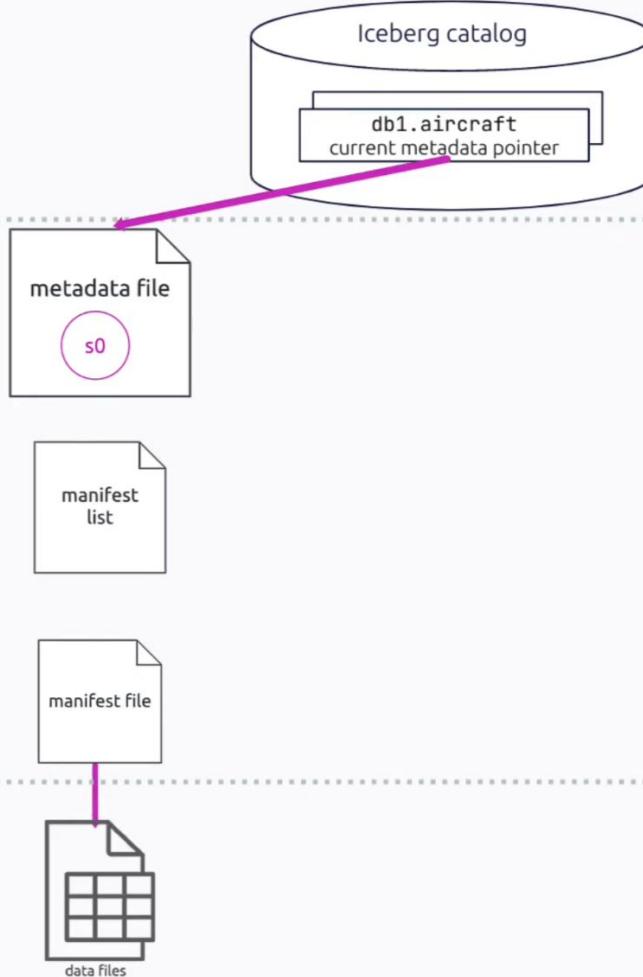
[Objects](#) | [Properties](#)**Objects (2)** [Info](#)[!\[\]\(3b451835b5cf44dc087a11f8c88642da_img.jpg\) C](#) [Copy S3 URI](#) [Copy URL](#) [Download](#) [Open](#) [Delete](#) [Actions ▾](#) [Create folder](#)[!\[\]\(0e60d2d9b679b4cf53dbe1e685ee345d_img.jpg\) Upload](#)

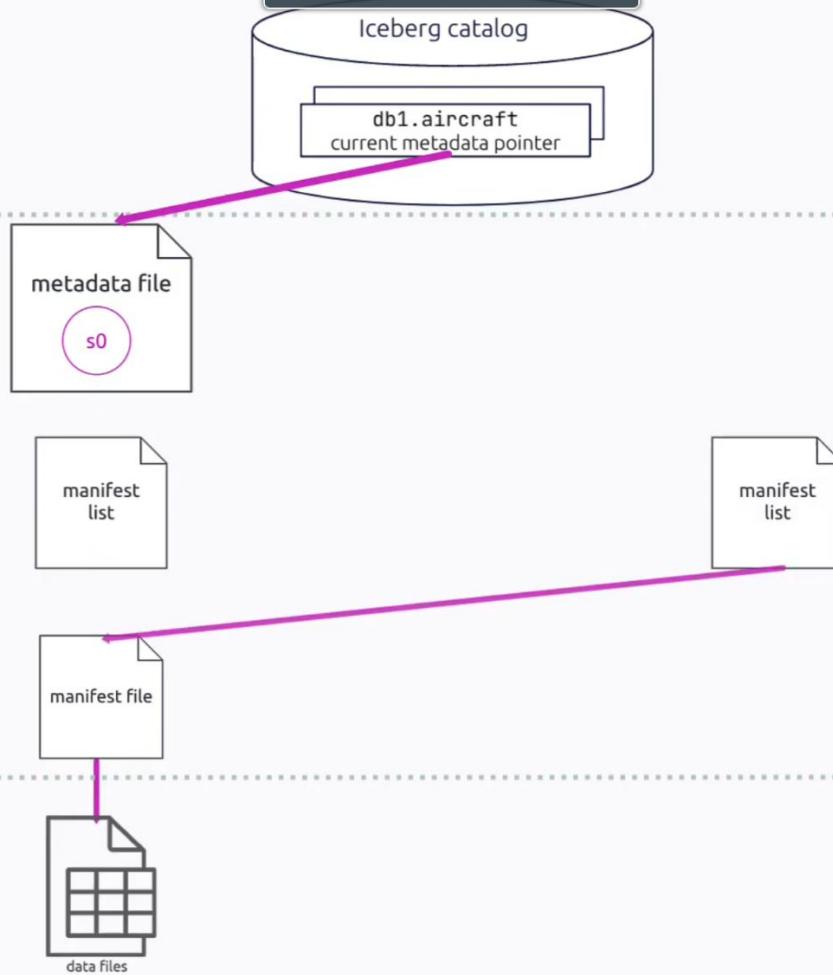
Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

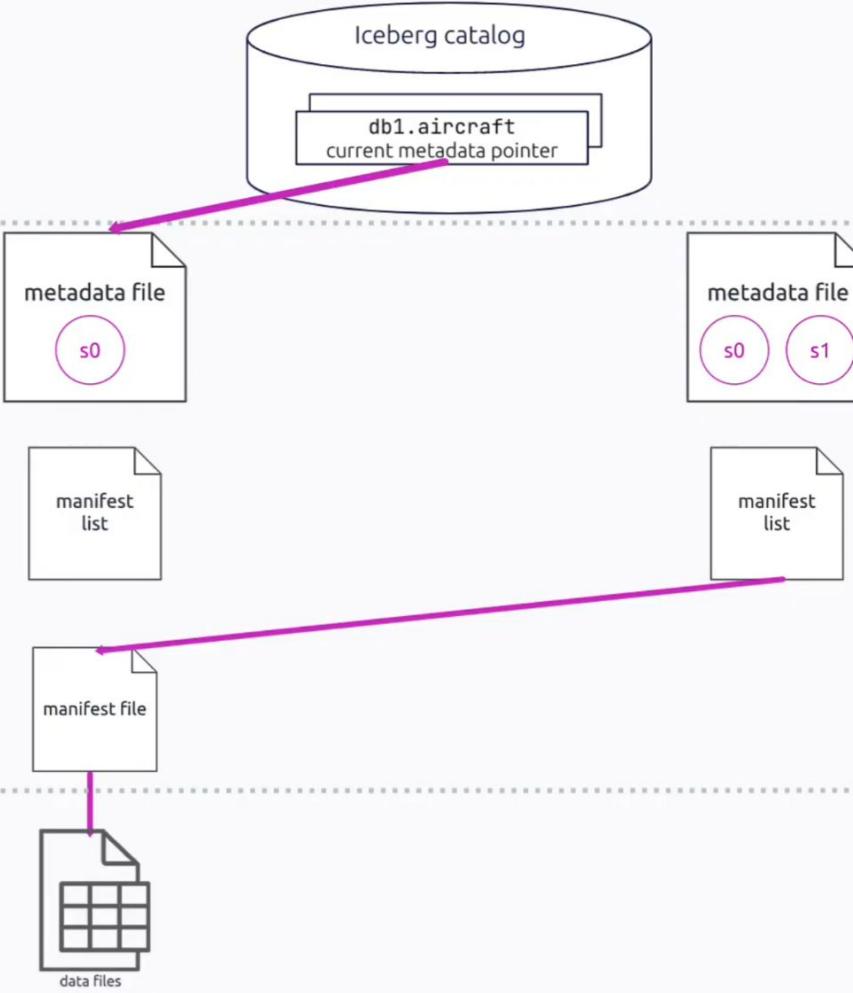
<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class	<input type="checkbox"/>
<input type="checkbox"/>	00000-e6b05edc-87c5-44c0-91ed-504767cd1107.metadata.json	json	May 28, 2024, 15:06:23 (UTC-04:00)	2.1 KB	Standard	
<input type="checkbox"/>	snap-425416382669527773-1-189af56b-cb7a-40b0-a0d1-6a94c3796ca0.avro	avro	May 28, 2024, 15:06:23 (UTC-04:00)	4.0 KB	Standard	

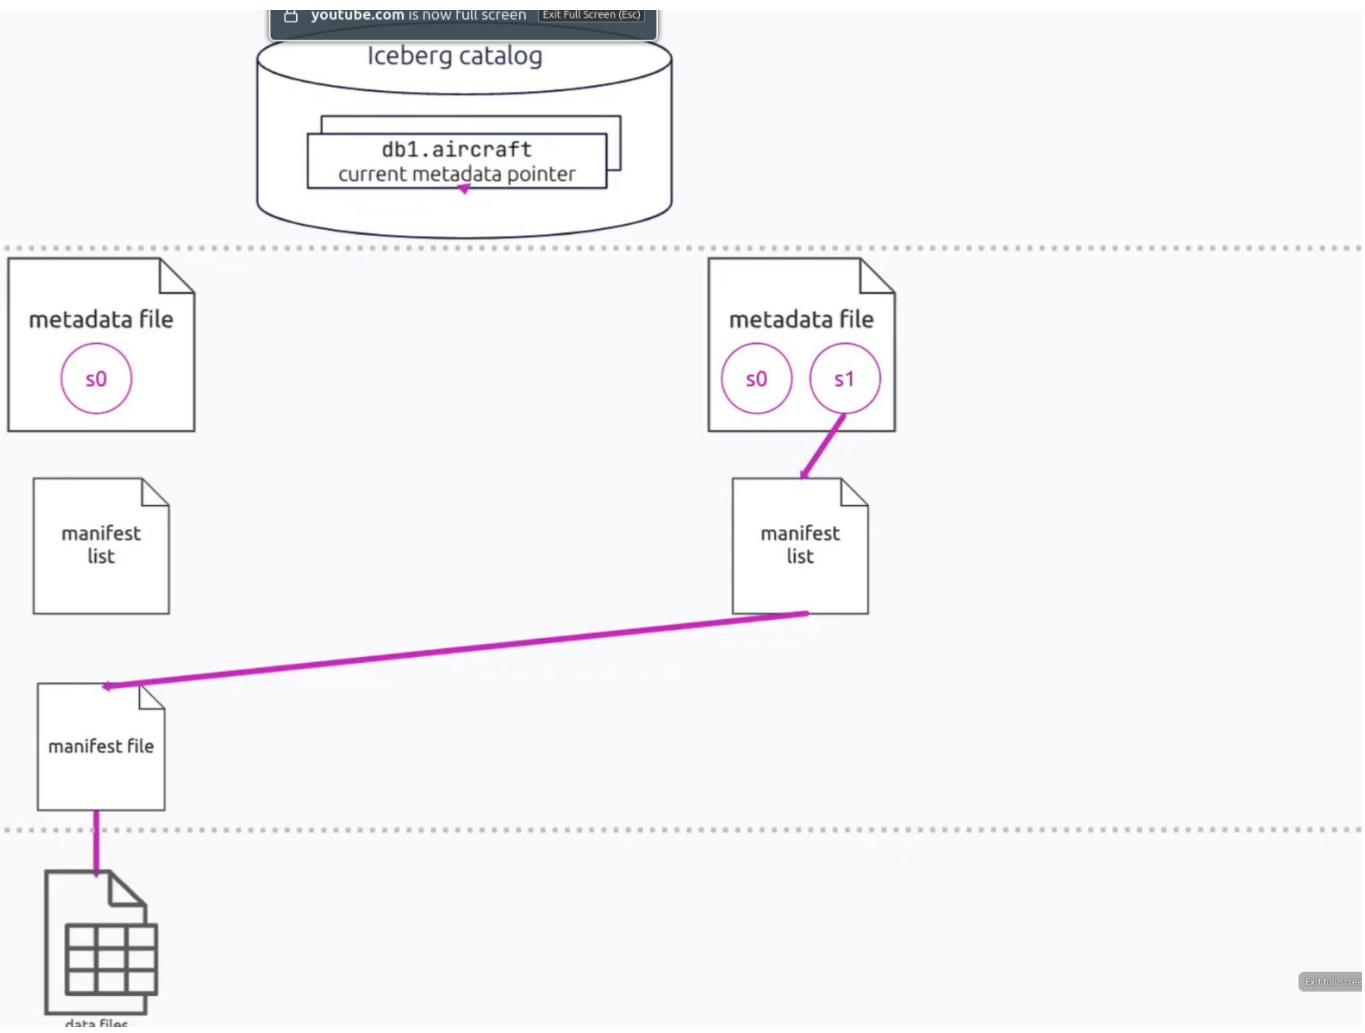
```
INSERT INTO
  aircraft (tail_number, description, class, year)
VALUES
  ('N535NA', 'NASA', 'Helicopter', 1969),
  ('N611TV', 'COOL', 'Jet', 1983);
```

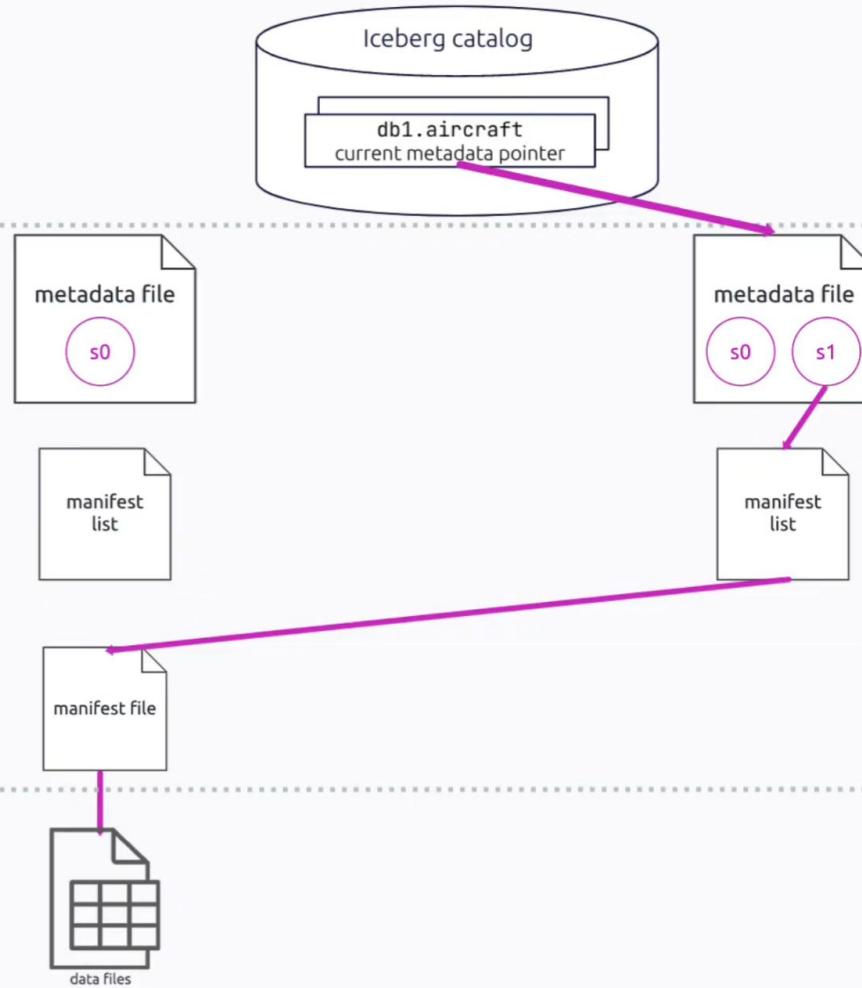












Objects (2) [Info](#)[Copy S3 URI](#)[Copy URL](#)[Download](#)[Open](#)[Delete](#)[Actions ▾](#)[Create folder](#)[Upload](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

 Find objects by prefix[<>](#) 1 [a](#) [g](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class	▼
<input type="checkbox"/>	data/	Folder	-	-	-	▼
<input type="checkbox"/>	metadata/	Folder	-	-	-	▼

Objects

Properties

Objects (1) Info

Copy S3 URI

Copy URL

Download

Open

Delete

Actions ▾

Create folder

Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

< 1 >

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class	
<input type="checkbox"/>	2024_0529_152110_2198		May 29, 2024, 11:22:37 (UTC-04:00)	658.0 B	Standard	
<input type="checkbox"/>	2_9r62a-588d9021- 5556-437f-a2f0- d0f00600f748.parquet	parquet				

Objects (7) Info

[Copy S3 URI](#) [Copy URL](#) [Download](#) [Open](#) [Delete](#) [Actions ▾](#) [Create folder](#)

[Upload](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

<input type="text"/> Find objects by prefix								1		
<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class	Actions				
<input type="checkbox"/>	20240529_152110_2198									
<input type="checkbox"/>	2_9r62a-740c70f1-774b-4b49-9067-03e1771f37a8.stats	stats	May 29, 2024, 11:22:38 (UTC-04:00)	1017.0 B	Standard					
<input type="checkbox"/>	29fa1a6e-e0aa-4aab-b9c3-58fa79dcfa44-m0.avro	avro	May 29, 2024, 11:22:38 (UTC-04:00)	6.7 KB	Standard					
<input type="checkbox"/>	snap-2724582809466504793-1-29fa1a6e-e0aa-4aab-b9c3-58fa79dcfa44.avro	avro	May 29, 2024, 11:22:38 (UTC-04:00)	4.2 KB	Standard					
<input type="checkbox"/>	snap-425416382669527773-1-189af56b-cb7a-40b0-a0d1-6a94c3796ca0.avro	avro	May 28, 2024, 15:06:23 (UTC-04:00)	4.0 KB	Standard					



Objects (7) Info

[!\[\]\(38c152d00a877a70f84c28a2599081db_img.jpg\) Create](#) [!\[\]\(cc18d2a09bb0deac541c36dd4dbf0dce_img.jpg\) Copy S3 URI](#) [!\[\]\(2b07aeab3b7baae2dd621042cb8ab4a4_img.jpg\) Copy URL](#) [!\[\]\(e323931596f077e429baf89af6d9585f_img.jpg\) Download](#) [!\[\]\(cb3e30ad64440eb02d4c60e7b1c92dd8_img.jpg\) Open](#) [!\[\]\(87d8cfcc36062a553146b6fc39e8524b_img.jpg\) Delete](#) [Actions ▾](#) [Create folder](#)

[!\[\]\(130eb552b7f729639752562c836f911d_img.jpg\) Upload](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

 Find objects by prefix



< 1 >



<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	20240529_152110_2198				
<input type="checkbox"/>	2_9r62a-740c70f1-774b-4b49-9067-03e1771f37a8.stats	stats	May 29, 2024, 11:22:38 (UTC-04:00)	1017.0 B	Standard
<input type="checkbox"/>	29fa1a6e-e0aa-4aab-b9c3-58fa79dcfa44-m0.avro	avro	May 29, 2024, 11:22:38 (UTC-04:00)	6.7 KB	Standard
<input type="checkbox"/>	snap-2724582809466504793-1-29fa1a6e-e0aa-4aab-b9c3-58fa79dcfa44.avro	avro	May 29, 2024, 11:22:38 (UTC-04:00)	4.2 KB	Standard
<input type="checkbox"/>	snap-425416382669527773-1-189af56b-cb7a-40b0-a0d1-6a94c3796ca0.avro	avro	May 28, 2024, 15:06:23 (UTC-04:00)	4.0 KB	Standard

Apache iceberg advantages

1) Reliable schema evolution

- add/drop/rename columns safely
- type promotion (where supported)

2) ACID-like consistency for the lake

With “just Parquet”, you typically don’t have atomic commits: readers can see partial writes, and concurrent writers can stomp on each other.
Iceberg uses **atomic snapshots**:

- writes commit as a new snapshot or not at all
- readers always see a consistent version
- concurrent writes are coordinated via metadata + catalog

3) Time travel and reproducibility

Iceberg keeps table history (snapshots), so you can:

- query “as of” a timestamp/snapshot
- rollback to a previous version
- reproduce ML features / reports exactly

4) Upserts, deletes, and MERGE that actually work

Parquet files are immutable; “updates” usually mean rewriting partitions and hoping you don’t break consistency.

Iceberg supports:

- row-level deletes (delete files / position deletes)
- MERGE INTO / UPDATE / DELETE patterns (engine support varies but the table format enables it)

5) Faster reads via metadata-driven pruning

Parquet has row-group stats, but without a table layer you still often end up listing tons of files and guessing what to scan.

Iceberg maintains metadata that enables:

- **partition pruning** without directory conventions
- **file pruning** via stats (min/max, null counts, etc.)
- less expensive “list” operations on object stores

8) Multi-engine interoperability (more predictable)

A Parquet directory structure might “work” in Spark but not in Trino/Presto the same way, especially with schema evolution and partition inference.

Iceberg’s spec makes behavior more consistent across engines that support it (Spark, Flink, Trino, Dremio, Athena, etc.).

When “pure Parquet” can be enough

- append-only pipelines with no updates/deletes
- single writer, low concurrency
- you don't need time travel, rollback, or evolving partition strategies